

COMP 4384 Software Security

Module 5: *Integer Overflow Attacks*

Ahmed Tamrawi



atamrawi



atamrawi.github.io



ahmedtamrawi@gmail.com

Acknowledgment Notice

Part of the slides are based on content from CMSC414 course by **Dave Levin** and **Niall Cooling**'s blog "When integers go bad" (<https://blog.feabhas.com/2014/10/vulnerabilities-in-c-when-integers-go-bad/>) and "Basic Integer Overflows" by **Phrack magazine** (<http://phrack.org/issues/60/10.html>)

What does the program print?

```
1 public class JavaPuzzle {  
2     public static void main(String[] args) {  
3         final long MICROS_PER_DAY = 24 * 60 * 60 * 1000 * 1000;  
4         final long MILLIS_PER_DAY = 24 * 60 * 60 * 1000;  
5         System.out.println(MICROS_PER_DAY / MILLIS_PER_DAY);  
6     }  
7 }
```

What does the program print?

```
1 public class JavaPuzzle {  
2     public static void main(String[] args) {  
3         final long MICROS_PER_DAY = 24 * 60 * 60 * 1000 * 1000;  
4         final long MILLIS_PER_DAY = 24 * 60 * 60 * 1000;  
5         System.out.println(MICROS_PER_DAY / MILLIS_PER_DAY);  
6     }  
7 }
```


It prints “5”!

What does the program print?

```
1 public class JavaPuzzle {  
2     public static void main(String[] args) {  
3         final long MICROS_PER_DAY = 24L * 60 * 60 * 1000 * 1000;  
4         final long MILLIS_PER_DAY = 24L * 60 * 60 * 1000;  
5         System.out.println(MICROS_PER_DAY / MILLIS_PER_DAY);  
6     }  
7 }
```

Takeaway 1 *When working with large numbers, watch out for **overflow**—it's a **silent killer**!*





What's Wrong with this Code?

```
void vulnerable()
{
    size_t len;
    char *buf;


    len = read_int_from_network();
    buf = malloc(len + 5);
    read(fd, buf, len);
    ...
}
```

What's Wrong with this Code?

```
void vulnerable()
{
    size_t len;
    char *buf;
    HUGE
    len = read_int_from_network();
    buf = malloc(len + 5); Wrap-around
    read(fd, buf, len);
    ...
}
```

Takeaway 2 *You have to know the semantics of your programming language to avoid these errors.*

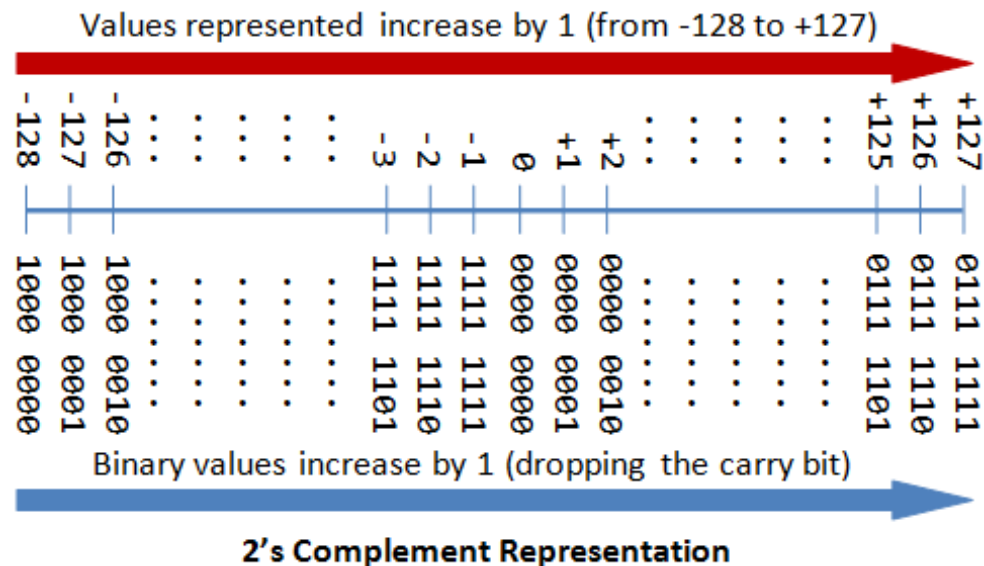
Integer Overflow Prevalence



<http://web.nvd.nist.gov/view/vuln/statistics>


Integers

- All built-in integral types (char, short, int, long, etc.) have a **limited capacity** because they are represented with a **fixed number of bits**.
- In most 32-bit architectures, signed integers (those that can be either positive or negative) are expressed in what is known as **two's complement notation**.



Integers

- Unlike integers in mathematics, program variables have a fixed range and “**wrap around**” when they go above their **maximum** value or below their **minimum** value; **a very large positive number becomes a very large negative number, or vice versa.**



Action: ADD 1

Bin: 0001

Hex: 1

Unsigned: 1


Signed: 1

Zero: 0

Carry: 0

Sign: 0

Overflow: 0



Type	Size(byte)	Symbol	Range
char/signed char	1	%c	-128 to 127
unsigned char	1	%c	0 to 255
Short int /signed short int	1	%d	-128 to 127
unsigned int	1	%u	0 to 255
int /signed int	2	%d	-32768 to 32767
unsigned int	2	%u	0 to 65535
long int/signed long int	4	%ld	-2147483648 to 2147483647
unsigned long int	4	%lu	0 to 4294967295
float	4	%f	-3.4e38 to 3.4e38
double	8	%lf	-3.4e308 to 3.4e308
long double	10	%lf	-1.7e4982 to 1.7e4982

What are the potential underlying problems of **fixed-sized representation** of numbers?

- Arithmetic Overflow
- Arithmetic Underflow
- Promotion/extension
- Demotion/narrowing
- Sign conversion

Arithmetic Overflow

- When an attacker can take advantage of this behavior, the program is said to contain an **integer overflow vulnerability**.
- Integer overflow can lead to any number of problems, but in C and C++, an **integer overflow** is most frequently used as a **lead-in to a buffer overflow exploit**.
 - The buffer overflow might occur when the wrapped-around variable is used to **allocate memory, bound a string operation, or index into a buffer**.
- Integer overflow can also occur in Java, but because Java enforces **memory safety properties**, integer overflow is not as easy to exploit.

Example 1: *Unsigned Overflow Vulnerability*

```
#include <stdio.h>

int main(void) {
    unsigned short a = 65000;
    unsigned short b = 540;
    unsigned short c = 0;

    c = a + b;
    printf("Result is %u + %u = %u\n", a, b, c);
    return 0;
}
```

overflow.c

Example 1: *Unsigned Overflow Vulnerability*

```
#include <stdio.h>

int main(void) {
    unsigned short a = 65000;
    unsigned short b = 540;
    unsigned short c = 0;

    c = a + b;
    printf("Result is %u + %u = %u\n", a, b, c);
    return 0;
}
```

```
local-admins-MacBook-Pro:module-05 ahmedtamrawi$ gcc overflow.c -o overflow
local-admins-MacBook-Pro:module-05 ahmedtamrawi$ ./overflow
Result is 65000 + 540 = 4
```

```
65000 => 0xfde8 => b'1111 1101 1110 1000
    540 => 0x021c => b'0000 0010 0001 1100
                        b'1 0000 0000 0000 0100
```


Example 2: *Arithmetic Underflow Vulnerability*

```
#include <stdio.h>
#include <limits.h>
#include <assert.h>

int main(void) {
    assert(sizeof(short)==2);

    unsigned short us = 0;
    short ss = SHRT_MIN; // -32768

    us -= 1;
    ss -= 1;

    printf("%u %d\n", us, ss);
    return 0;
}
```

underflow.c

Example 2: *Arithmetic Underflow Vulnerability*

```
#include <stdio.h>
#include <limits.h>
#include <assert.h>

int main(void) {
    assert(sizeof(short)==2);

    unsigned short us = 0;
    short ss = SHRT_MIN; // -32768


    us -= 1;
    ss -= 1;

    printf("%u %d\n", us, ss);
    return 0;
}
```

```
local-admins-MacBook-Pro:module-05 ahmedtamrawi$ gcc underflow.c -o underflow
local-admins-MacBook-Pro:module-05 ahmedtamrawi$ ./underflow
65535 32767
```

Integer Promotion/Extension

- **Type promotion** occurs when we convert from a **small sized integer** to a **larger one**, e.g. from **short** to **int**.
- For example, when a signed integer is converted from a smaller number of bits to a larger number of bits, the extra bits are *filled in so that the new number retains the same sign*.
 - Negative number casted to signed larger data type, its signed value will remain the same. (1000 -> 1111 1000)
 - Negative number casted to unsigned larger data type will increase significantly because its most significant bits will be set. (1000 -> 1111 1000)



Example 1: *Integer Promotion/Extension*

```
#include <stdio.h>
#include <limits.h>
#include <assert.h>

int main(void) {
    assert(sizeof(short)==2);

    short ss = SHRT_MIN;
    int si = ss;

    printf("%d %d\n", ss, si);
    printf("%x %x\n", ss, si);

    return 0;
}
```

signedPromotion.c

Example 1: *Integer Promotion/Extension*

```
#include <stdio.h>
#include <limits.h>
#include <assert.h>

int main(void) {
    assert(sizeof(short)==2);

    short ss = SHRT_MIN;
    int si = ss;

    printf("%d %d\n", ss, si);
    printf("%x %x\n", ss, si);

    return 0;
}
```

```
local-admins-MacBook-Pro:module-05 ahmedtamrawi$ gcc signedPromotion.c -o signedPromotion
local-admins-MacBook-Pro:module-05 ahmedtamrawi$ ./signedPromotion
-32768 -32768
8000 ffff8000
```

Example 2: *Integer Promotion/Extension*

```
#include <stdio.h>
#include <limits.h>
#include <assert.h>

int main(void) {
    assert(sizeof(short)==2);

    short ss = SHRT_MIN;
    unsigned int si = ss;

    printf("%d %u\n", ss, si);
    printf("%x %x\n", ss, si);

    return 0;
}
```

unsignedPromotion.c

Example 2: *Integer Promotion/Extension*

```
#include <stdio.h>
#include <limits.h>
#include <assert.h>

int main(void) {
    assert(sizeof(short)==2);

    short ss = SHRT_MIN;
    unsigned int si = ss;

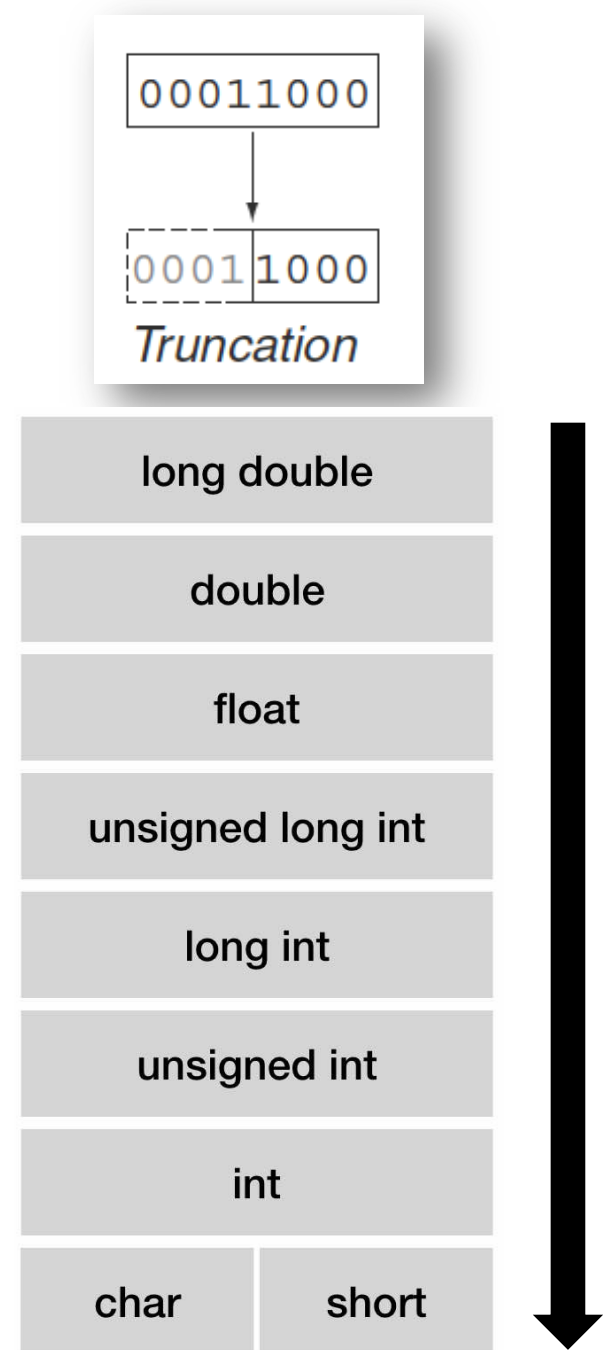
    printf("%d %u\n", ss, si);
    printf("%x %x\n", ss, si);

    return 0;
}
```

```
local-admins-MacBook-Pro:module-05 ahmedtamrawi$ gcc unsignedPromotion.c -o unsignedPromotion
local-admins-MacBook-Pro:module-05 ahmedtamrawi$ ./unsignedPromotion
-32768 4294934528
ffff8000 ffff8000
```

Integer Demotion/Narrowing

- **Integer truncation** errors occur when an integer data type with a larger number of bits is converted to a data type with fewer bits.
- Narrowing occurs through *truncating the bits to the target type's size*.
 - For example, going from `int` to `short` will result in the bottom 16-bits of the 32-bit `int` being copied to the `short`.
 - For **unsigned numbers**, this may result in a **loss of information** (i.e. large numbers being truncated to small numbers).
 - For **signed numbers**, narrowing can result in unexpected change of sign.



Example 1: *Integer Demotion/Narrowing*

```
#include <stdio.h>
#include <limits.h>
#include <assert.h>

#define MAGIC_NUMBER 0xFFFF7F8F

int main(void) {
    assert(sizeof(short)==2);

    unsigned int ui = MAGIC_NUMBER;
    unsigned short us = ui;
    unsigned char uc = us;

    int si = MAGIC_NUMBER;
    short ss = si;
    signed char sc = ss;

    printf("%10u %5hu %4hhu\n", ui, us, uc);
    printf("%10x %5hx %4hhx\n", ui, us, uc);
    printf("%10d %5hd %4hhd\n", si, ss, sc);
    printf("%10x %5hx %4hhx\n", si, ss, sc);
    return 0;
}
```

```

#include <stdio.h>
#include <limits.h>
#include <assert.h>

#define MAGIC_NUMBER 0xFFFF7F8F

int main(void) {
    assert(sizeof(short)==2);

    unsigned int ui = MAGIC_NUMBER;
    unsigned short us = ui;
    unsigned char uc = us;

    int si = MAGIC_NUMBER;
    short ss = si;
    signed char sc = ss;

    printf("%10u %5hu %4hhu\n", ui, us, uc);
    printf("%10x %5hx %4hhx\n", ui, us, uc);
    printf("%10d %5hd %4hhd\n", si, ss, sc);
    printf("%10x %5hx %4hhx\n", si, ss, sc);
    return 0;
}

```

```



local-admins-MacBook-Pro:module-05 ahmedtamrawi$ gcc narrowing.c -o narrowing
local-admins-MacBook-Pro:module-05 ahmedtamrawi$ ./narrowing
4294934415 32655 143
ffff7f8f 7f8f 8f
-32881 32655 -113
ffff7f8f 7f8f 8f

```

Note the change in sign for short

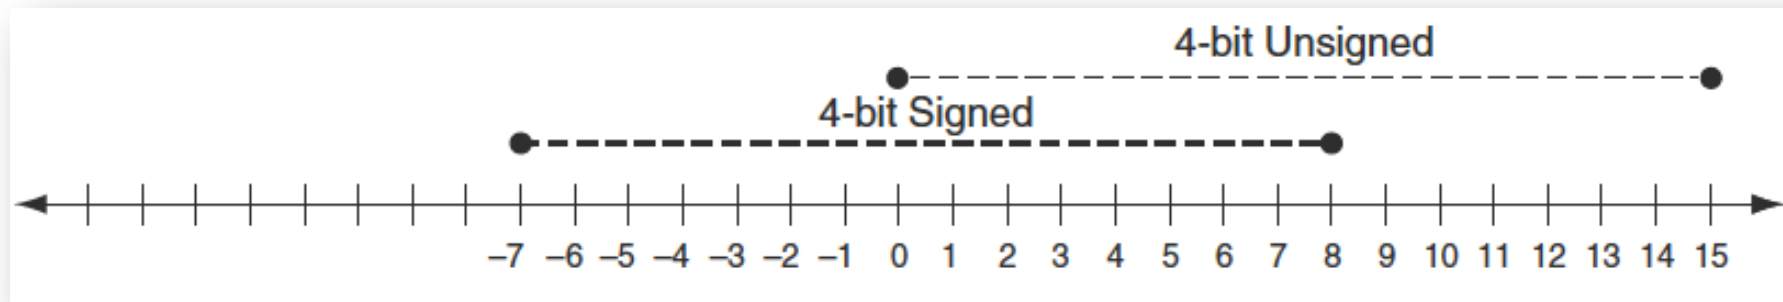
Sign Conversion

- Both **signed** and **unsigned** data types are capable of representing the same number of values because they have the same number of bits available to them.
 - However there is only **partial overlap** between the range of numbers that the two types can express.



Sign Conversion

- The result of this **partial overlap** is that some values can be converted from an *unsigned data type* to a *signed data type* and vice versa without a **change** in meaning, while **others cannot**.
- Intuitively, this is the case for signed-to-unsigned conversions **because a negative value cannot be represented as an unsigned data type**.



Sign Conversion

- In the case of positive values, the problem is that the largest 50% of unsigned values require setting the high-order bit.
- The same bit pattern interpreted as a signed quantity will be negative.
 - If the most-significant-bit (MSB) is a zero (0) then there are no issues with the conversion in either direction.
 - If, however, the MSB is a 1 then a change in sign and value will occur.

$\boxed{1000}$	$= -8$	$\boxed{1000}$	$= 15$
<i>Signed</i>		<i>Unsigned</i>	

Example 1: *Sign Conversion*

```
#include <stdio.h>
#include <limits.h>
#include <assert.h>

int main(void) {
    assert(sizeof(short)==2);

    unsigned short us = 0x8080;
    short ss = us;

    printf("%6hu %6hd\n", us, ss);
    printf("%6hx %6hx\n", us, ss);
    return 0;
}
```

conversion.c

Example 1: *Sign Conversion*

```
#include <stdio.h>
#include <limits.h>
#include <assert.h>

int main(void) {
    assert(sizeof(short)==2);

    unsigned short us = 0x8080;
    short ss = us;

    printf("%6hu %6hd\n", us, ss);
    printf("%6hx %6hx\n", us, ss);
    return 0;
}
```

```
local-admins-MacBook-Pro:module-05 ahmedtamrawi$ gcc conversion.c -o conversion
local-admins-MacBook-Pro:module-05 ahmedtamrawi$ ./conversion
 32896 -32640
  8080   8080
```

Arithmetic Conversion/Promotion

- So far, we have mostly focused on types of the same size (e.g. `short` and `unsigned short`), but if we have arithmetic or logic operations a pattern called the *usual arithmetic conversions* are applied.
- This means, that for **arithmetic and logic operations**, integer types shorter than an `int` are **promoted to an `int` for the operation**.
 - The promotions can sometimes lead to **unexpected consequences**, such as signed values being interpreted as unsigned and vice versa.

Example 1: *Arithmetic Conversion/Promotion*

```
#include <stdio.h>
#include <limits.h>
#include <assert.h>

int main(void) {
    assert(sizeof(unsigned char)==1);

    unsigned char uc1 = 0xff;
    unsigned char uc2 = 0;

    if(~uc1 == uc2) {
        printf("%hhx == %hhx\n", ~uc1, uc2);
    } else {
        printf("%hhx != %hhx\n", ~uc1, uc2);
    }
    return 0;
}
```

promotion.c

```

#include <stdio.h>
#include <limits.h>
#include <assert.h>

int main(void) {
    assert(sizeof(unsigned char)==1);

    unsigned char uc1 = 0xff;
    unsigned char uc2 = 0;

    if(~uc1 == uc2) {
        printf("%hhx == %hhx\n", ~uc1, uc2);
    } else {
        printf("%hhx != %hhx\n", ~uc1, uc2);
    }
    return 0;
}

```

```

local-admins-MacBook-Pro:module-05 ahmedtamrawi$ gcc promotion.c -o promotion
promotion.c:13:30: warning: format specifies type 'unsigned char' but the argument has type 'int' [-Wformat]
    printf("%hhx == %hhx\n", ~uc1, uc2);
                        ~~~~      ^~~~~
                        %X
promotion.c:15:30: warning: format specifies type 'unsigned char' but the argument has type 'int' [-Wformat]
    printf("%hhx != %hhx\n", ~uc1, uc2);
                        ~~~~      ^~~~~
                        %X
2 warnings generated.
local-admins-MacBook-Pro:module-05 ahmedtamrawi$ ./promotion
0 != 0

```

Example 2: *Arithmetic Conversion/Promotion*

As uc1 has been promoted to the unsigned integer 0x000000ff, when complimented it results in 0xffffffff00, as shown and thus not equal to zero.

```
#include <stdio.h>
#include <limits.h>
#include <assert.h>

int main(void) {
    assert(sizeof(unsigned char)==1);

    unsigned char uc1 = 0xff;
    unsigned char uc2 = 0;

    if(~uc1 == uc2) {
        printf("%08x == %08x\n", ~uc1, uc2);
    } else {
        printf("%08x != %08x\n", ~uc1, uc2);
    }
    return 0;
}
```

promotion2.c

```
local-admins-MacBook-Pro:module-05 ahmedtamrawi$ gcc promotion2.c -o promotion2
local-admins-MacBook-Pro:module-05 ahmedtamrawi$ ./promotion2
ffffffff00 != 00000000
```

INT_MIN

- There is one other anomaly to be aware of based around INT_MIN. When using 2's complement the number range of an integer is not symmetrical, i.e. the range is:
$$-2147483648 \dots 2147483647$$
- All negative values, apart from INT_MIN, have a positive representation. Unfortunately we cannot represent -2147483648 as a positive signed number.
- This leads to the strange behavior that the absolute of **INT_MIN** and **-INT_MIN** both are likely to yield **INT_MIN**.

Example 1: INT_MIN

```
#include <stdio.h>
#include <limits.h>
#include <assert.h>
#include <stdlib.h>

int main(void) {
    assert(sizeof(int)==4);

    int intMin = INT_MIN;

    printf("%d %d %d\n", intMin, abs(intMin), -intMin);
    return 0;
}
```

intMin.c

```
local-admins-MacBook-Pro:module-05 ahmedtamrawi$ gcc intMin.c -o intMin
local-admins-MacBook-Pro:module-05 ahmedtamrawi$ ./intMin
-2147483648 -2147483648 -2147483648
```

Example 1: *Exposing Integer Overflow Vulnerability for Privilege Escalation Attack*

- Suppose a network service keeps track of the number of connections it has received since it has started, and only grants access to the **first five users**.

```
#include <stdio.h>

int main(int argc, char * argv[])
{
    unsigned int connections = 0;
    // Insert network code here
    // ...
    // ...
    // Does nothing to check overflow conditions
    connections++;
    if(connections < 5)
        grant_access();
    else
        deny_access();
    return 1;
}
```

Example 1: *Exposing Integer Overflow Vulnerability for Privilege Escalation Attack*

- An attacker could compromise the above system by making a huge number of connections until the connections **counter overflows** and wraps around to zero.
- At this point, the attacker will be authenticated to the system, which is clearly an undesirable outcome.

```
#include <stdio.h>

int main(int argc, char * argv[])
{
    unsigned int connections = 0;
    // Insert network code here
    // ...
    // ...
    // Does nothing to check overflow conditions
    connections++;
    if(connections < 5)
        grant_access();
    else
        deny_access();
    return 1;
}
```

Example 1: *Exposing Integer Overflow Vulnerability for Privilege Escalation Attack*

```
#include <stdio.h>

int main(int argc, char * argv[])
{
    unsigned int connections = 0;
    // Insert network code here
    // ...
    // ...
    // Does nothing to check overflow conditions
    connections++;
    if(connections < 5)
        grant_access();
    else
        deny_access();
    return 1;
}
```



```
#include <stdio.h>

int main(int argc, char * argv[])
{
    unsigned int connections = 0;
    // Insert network code here
    // ...
    // ...
    // Prevents overflow conditions
    if(connections < 5)
        connections++;
    if(connections < 5)
        grant_access();
    else
        deny_access();
    return 1;
}
```


Example 2: *Integer Underflow Vulnerability*

- The most common root problem using integer-based attacks is where the implementation of an algorithm has mixed signed and unsigned values.
- Good targets are where standard library functions, such as `malloc` or `memcpy` have been used, as in both cases they take parameters of type `size_t` (unsigned integer data type).

```
int copySize;  
// do work, copySize calculated..  
if (copySize > MAX_BUF_SZ) {  
    return -1;  
}  
memcpy(&d, &s, copySize*sizeof(type));
```

Defense Against the Dark Arts

- In short, it can be very **difficult** to protect ourselves against building programs which accidentally or deliberately use the undefined or implementation defined integer behavior. Nevertheless, there are several things we can do:
 - Education
 - Use your compiler flags
 - Follow a Security based coding standard
 - Enforce the Coding Standard using a Static Analysis (SA) Tool